


MICHAL ROSTECKI

Rust Engineer at Exein

Enhancing Rust with BTF Debug Format Support



► INTRODUCTION

- What is eBPF?
 - How it differs from kernel modules?
 - What is BTF? How is it produced?
- **You can write eBPF in Rust!** But there is still some work to be done to meet 100% feature parity.

► INTRODUCTION

- There are two ways of extending the Linux kernel
 - Kernel modules
 - eBPF programs



The difference

Kernel modules vs eBPF

- **Kernel modules**
 - Have full access to kernel internals.
 - Need kernel sources for building.
 - It's encouraged to submit them upstream, to the main kernel repo.
 - **Can crash the kernel.**
 - **Great for device drivers, filesystems and extending kernel's functionality directly.**

The difference

Kernel modules vs eBPF




- **eBPF**
 - Built independently from the kernel. You just need a compiler supporting BPF target.
 - **Cannot crash the kernel.**
 - **Great for third-party tools which hook into Linux internals - firewalls, tracers, security monitors.**

Tell me more


How eBPF works?

- **It's a virtual machine**
 - It comes with its own architecture and assembly.
 - Forces programs to be small, on purpose.
 - That's why it can't crash the kernel.
 - Compilers (LLVM and GCC) support BPF target.

- **Kernel modules vs BPF** C vs Rust



LIBBPF



Writing BPF in Rust

Aya



- **Rust library for writing eBPF**
 - <https://aya-rs.dev>
 - Rustup and cargo is all you need!

Security observability framework
for IoT devices (but not only!)

Pulsar



- **Security monitor, using Rust and Aya**
 - <https://pulsar.sh/>
 - Lightweight.
 - Designed mostly for IoT devices. Has support for ARM and RISC-V.
 - But useful on x86 as well.

- **Using eBPF for network tracing** XDP



- **Using eBPF for network filtering** XDP








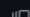






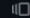


Here be dragons

What is Debug Info?

- Metadata that is generated by compilers and stored alongside the binary.
- Basically what debuggers use to figure out which line of code are you in.
- **But there is also an another use case!**

- Let's look at some kernel type to trace Whoops, we have a problem!

18 years ago	Linux-2.6.12-rc2		746	<code>struct task_struct {</code>
7 years ago	 sched/core: Allow putting threa...		747	<code>#ifdef CONFIG_THREAD_INFO_IN_TASK</code>
			748	<code>/*</code>
			749	<code> * For reasons of header soup (see current_thread_info()), this</code>
			750	<code> * must be the first element of task_struct.</code>
			751	<code>*/</code>
7 years ago	sched/headers: Clean up <linux...		752	<code>struct thread_info thread_info;</code>
7 years ago	 sched/core: Allow putting threa...		753	<code>#endif</code>
2 years ago	sched: Change task_struct::state		754	<code>unsigned int __state;</code>
6 years ago	 task_struct: Allow randomized I...		755	
2 years ago	sched/wakeup: Prepare for RT s...		756	<code>/* saved state for "spinlock sleepers" */</code>
			757	<code>unsigned int saved_state;</code>
			758	
6 years ago	 task_struct: Allow randomized I...		759	<code>/*</code>
			760	<code> * This begins the randomizable portion of task_struct. Only</code>
			761	<code> * scheduling-critical items should be added above here.</code>
			762	<code>*/</code>
			763	<code>randomized_struct_fields_start</code>
			764	
7 years ago	sched/headers: Clean up <linux...		765	<code>void *stack;</code>
4 years ago	 sched/core: Convert task_struct...		766	<code>refcount_t usage;</code>
7 years ago	sched/headers: Clean up <linux...		767	<code>/* Per task flags (PF_*), defined further below: */</code>
			768	<code>unsigned int flags;</code>
			769	<code>unsigned int ptrace;</code>

How to support many kernel versions?

- Different kernel versions can have different definitions of types you want to inspect.
- We need some mechanism for handling these differences.
- **Guess what... debug info is helpful!**

The most popular format

DWARF

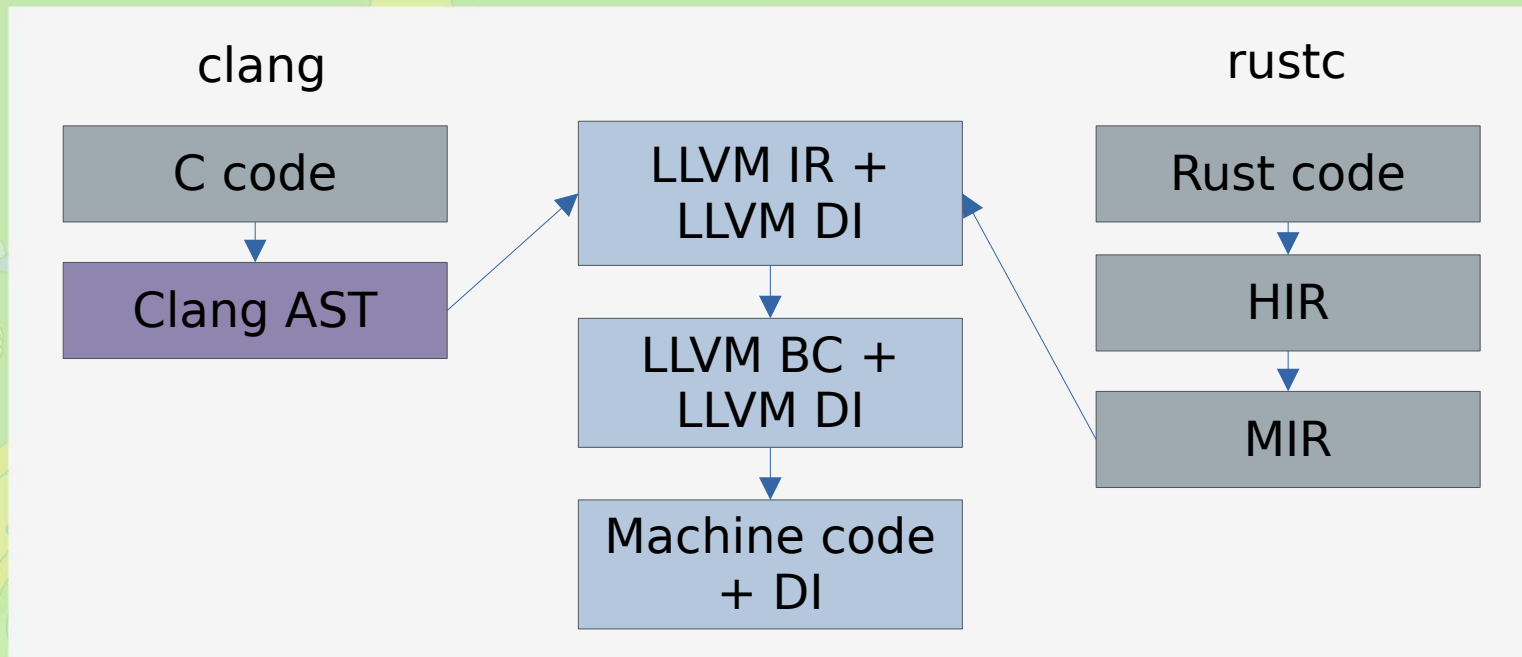
- The most popular debug information standard.
- Again, that's basically what your debugger uses!
- **It's quite big! It can take megabytes.**

Here it comes

BTF

- **BPF Type Format**
- Enables single compilation of eBPF programs for various kernel versions.
- Minified information about types and their fields with offsets.
- Enhanced debugging: Stack trace visibility when eBPF program load fails

- **LLVM** How does Debug Information get produced?



- **LLVM Debug Info** Example

```
!28223 = distinct !DICompositeType(tag: DW_TAG_structure_type, name:
"sockaddr_in6", file: !6, line: 22816, size: 224, elements: !28224)

!28224 = !{!28225, !28226, !28227, !28228, !28229}

!28225 = !DIDerivedType(tag: DW_TAG_member, name: "sin6_family", scope: !28223,
file: !6, line: 22817, baseType: !682, size: 16)

!28226 = !DIDerivedType(tag: DW_TAG_member, name: "sin6_port", scope: !28223, file:
!6, line: 22818, baseType: !990, size: 16, offset: 16)

!28227 = !DIDerivedType(tag: DW_TAG_member, name: "sin6_flowinfo", scope: !28223,
file: !6, line: 22819, baseType: !971, size: 32, offset: 32)

!28228 = !DIDerivedType(tag: DW_TAG_member, name: "sin6_addr", scope: !28223, file:
!6, line: 22820, baseType: !9016, size: 128, offset: 64)

!28229 = !DIDerivedType(tag: DW_TAG_member, name: "sin6_scope_id", scope: !28223,
file: !6, line: 22821, baseType: !899, size: 32, offset: 192)
```

- **BTF** Example

```
[200] STRUCT 'sockaddr_in6' size=28 vlen=5
      'sin6_family' type_id=19 bits_offset=0
      'sin6_port' type_id=24 bits_offset=16
      'sin6_flowinfo' type_id=14 bits_offset=32
      'sin6_addr' type_id=35 bits_offset=64
      'sin6_scope_id' type_id=15 bits_offset=192
```

How do they work?

- BPF programs contain the BTF for all types.
- Fields of these types are accessed by a compiler intrinsic – `preserve_access_index`.
- `preserve_access_index` gets compiled to a CO-RE relocation in BPF assembly.
- BPF VM knows what to do the relocation.

- **preserve_access_index** Intrinsic

```
declare <ret_type>
```

```
@llvm.preserve.union.access.index.p0s_union.anons.p0s_union.anons(  
    <type> base, i32 di_index)
```

```
declare <ret_type>
```

```
@llvm.preserve.struct.access.index.p0i8.p0s_struct.anon.0s(  
    <type> base, i32 gep_index, i32 di_index)
```

- **BTF relocations** How do they work?

Older kernel

MyStruct

u32 old_field_a

u32 old_field_b

Newer kernel

MyStruct

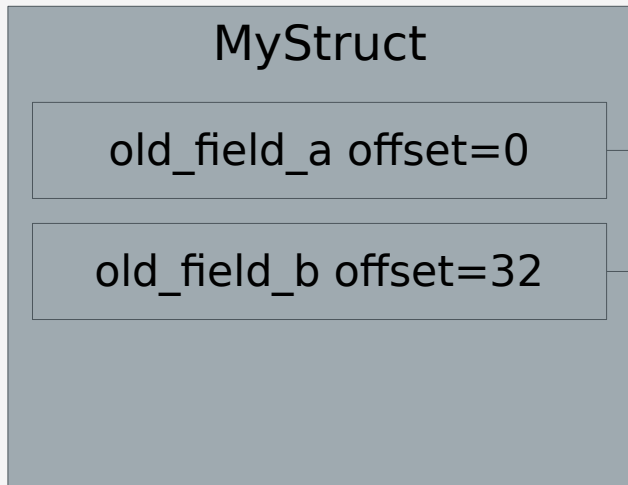
u32 old_field_a

u32 new_field_a

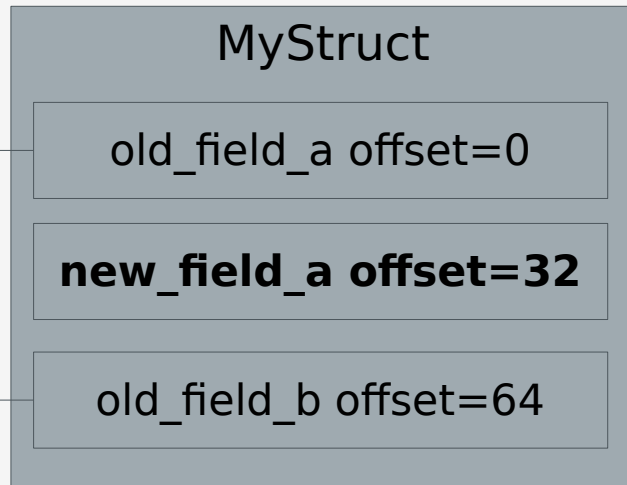
u32 old_field_b

- **BTF relocations** How do they work?

Program's BTF



Kernel's BTF



BTF - Aya's Feature Gap

- Rust doesn't support `preserve_access_index` intrinsic.
- In theory, emitting BTF should just work. **But it didn't work.**
 - BTF was made with assumptions about C types.
 - Initially, it was crashing the LLVM BPF backend (we fixed it).
 - Kernel assumes C types as well.

Be more precise, please

What does the kernel assume?

- **It gets annoyed by Rust-specific types**
 - Data-carrying enums.
 - Non-alphanumeric type names (e.g. types with generics).
- **It expects types Rust doesn't support**
 - BPF map types have to be anonymous structs.
 - Universal eBPF program compilation for all kernel versions
 - Simplified debugging with stack traces

How can we fix it

Stages of BTF support

- Emitting BTF
 - **1st stage: sanitizing LLVM Debug Info in bpf-linker**
 - 2nd stage: teaching kernel to support Rust types
- BTF relocations
 - `preserve_access_index` in Rust compiler

The 1st stage

Sanitizing LLVM DI

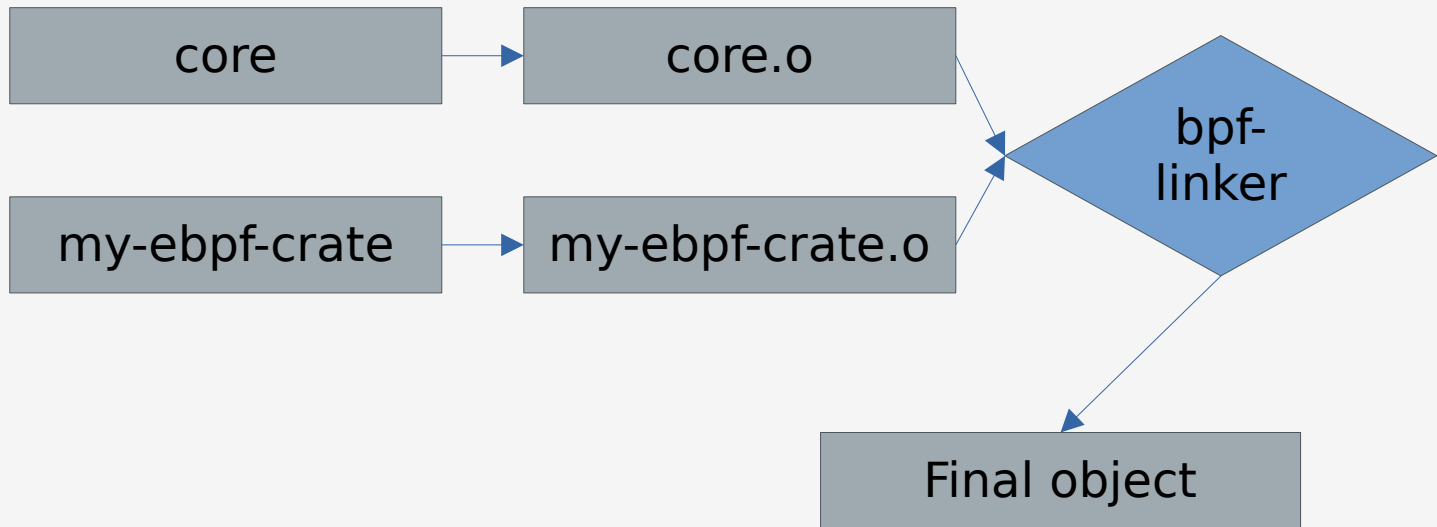
- **It can be done in bpf-linker.**
- To meet the expectations:
 - Remove children of data carrying enums.
 - Introduce a marker for anonymous types.
 - Sanitize names of all types.

But...

What is bpf-linker?

- It's a bitcode linker.
- Linking BPF in C is optional and mostly not done.
- **Linking in Rust is mandatory, everything is a crate.**
- Traditional linkers (e.g. lld) don't work for BPF.

- **How bpf-linker works?** Linking crates as bitcodes



A sneaky solution

Sanitizing LLVM DI

- Work in progress on feature/fix-di branch of github.com/aya-rs/bpf-linker
- Shoutout to everyone making it happen!
 - <https://github.com/davibe>
 - <https://github.com/qjerome>
 - <https://github.com/tamird>
- Before releasing, we want to
 - Provide test cases with the whole Rust types spectrum.
 - Statically link LLVM to bpf-linker.

Adding the BTF relocation support
in

Rust compiler


- `preserve_access_index` is an LLVM IR intrinsic.
- Works similar to GEP (`getElementPtr`) instruction.
- We need to add it to **`core::intrinsics`** and **`rustc_codegen_llvm`**.

Long-term, correct solution

Teaching the kernel to accept non-C types

- Probably would make sense to do when introducing BTF to Rust-for-Linux (**kernel modules need BTF for their types**).
- Would be nice to teach Rust to emit only BTF (other way than debug=2).
- Downside: will take years to be adopted.

Thank you for listening!



MICHAL ROSTECKI

michal@exein.io

