

AYA

Extending the Linux Kernel with eBPF and Rust



ABOUT ME

- Michal Rostecki
- **vadorovsky** @ Github, Discord, Twitter etc.
- Software Engineer @ Deepfence Inc
- Rustacean with Go, C and Python background

WHAT IS EBPf?

- extended Berkeley Packet Filter
- A virtual machine with its own instruction set which runs sandboxed programs.
- Originated in the Linux kernel.

EVENT-DRIVEN

- Triggered by events in the kernel.
- Receiving pointers to kernel or userspace memory.

WHAT EVENTS?

NETWORK TRAFFIC

- **XDP** - ingress traffic on NIC driver
- **Classifier** - both directions, attached to qdisc

FUNCTION CALLS

- Kprobes - kernel functions
- Uprobes - userspace functions

SECURITY-RELATED ACTIONS

- LSM hooks

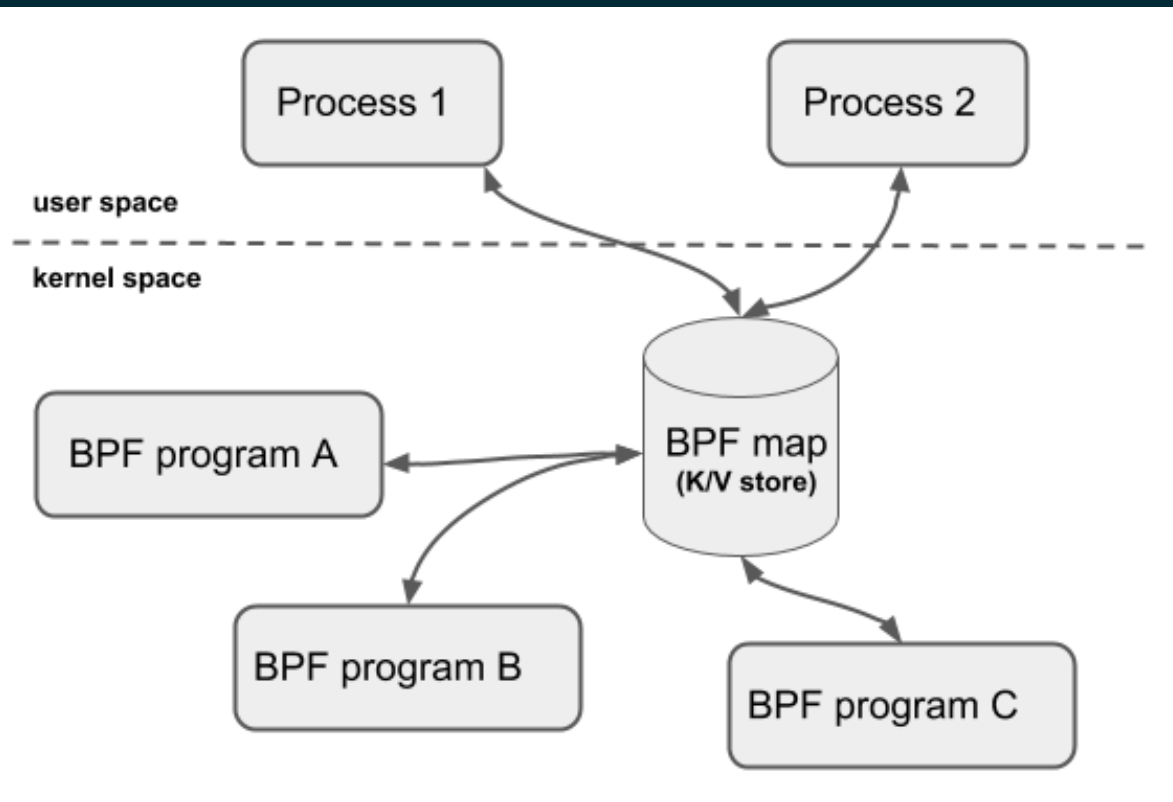
AND MANY MANY MORE...

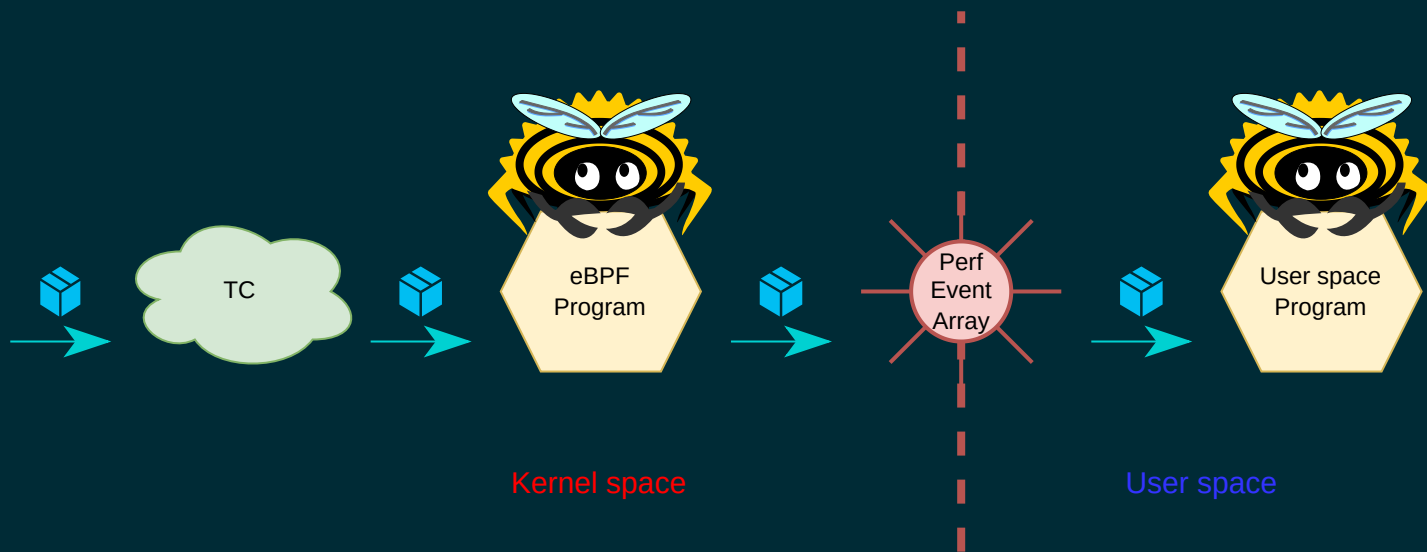
EBPF PROJECT LAYOUT

- **eBPF programs:** running in the kernel and reacting to events
- **userspace program:** loading eBPF programs to the kernel, interacting with them

BPF MAPS

- Storage for sharing data between eBPF programs (in kernel) and userspace.
- Different types:
 - hash maps
 - arrays
 - ring / perf buffers





EXAMPLES OF EBPf PROJECTS

- Cilium - Kubernetes networking, service mesh
- Tracee - application tracing tool
- Falco - security monitoring for containers

(those are in Go in C)

EXAMPLES OF EBPf PROJECTS USING AYA

- **bpfd** - daemon managing XDP programs for network filtering
- **Pulsar** - runtime observability tool for IoT

"OFFICIAL" WAYS TO USE EBPF

- libbpf
- libbpfgo
- libbpf-rs
- bcc

(With all of these, eBPF code is still in C)

AYA

Everything in Rust, even the eBPF programs!



WHY AYA?



RUST AND LINUX KERNEL

- There is [Rust for Linux](#).
- There is an ongoing race to rewrite parts of the kernel in eBPF as well.
- *Aya says: why not both?*

MEMORY SAFETY?

- Actually... Rust memory model doesn't matter that much in eBPF part.
- eBPF has access to kernel and userspace pointers, which are unsafe.
- Memory safety relies more on the **eBPF verifier**.

C AND TYPE SAFETY

```
SEC("xdp")
int incorrect_xdp(struct __sk_buff *skb) {
    return XDP_PASS;
}
```

We made a mistake! We should have used

```
struct xdp_md *ctx
```

instead of

```
struct __sk_buff *skb
```

```
$ clang -O2 -emit-llvm -c incorrect_xdp.c -o - | llc -march=bpf -  
$
```

But it compiles anyway, huh.

RUST AND TYPE SAFETY

```
#[xdp(name = "incorrect_xdp")]  
pub fn incorrect_xdp(ctx: SkBuffContext) -> u32 {  
    xdp_action::XDP_PASS  
}
```

It does not compile, nice.

ERROR HANDLING IN C

```
struct {  
    __uint(type, BPF_MAP_TYPE_HASH);  
    __uint(max_entries, 1024);  
    __type(key, pid_t);  
    __type(value, u32);  
} pids SEC(".maps");
```

```
SEC("fentry/kernel_clone")  
int BPF_PROG(kernel_clone, struct kernel_clone_args *args)  
{  
    /* Get the pid */  
    pid_t pid = bpf_get_current_pid_tgid() >> 32;  
    /* Save the pid in map */  
    u32 val = 0;  
    int err = bpf_map_update_elem(&pids, &pid, &val, 0);  
    if (err < 0)  
        return err;  
    return 0;  
}
```

ERROR HANDLING IN RUST

```
#[map(name = "pids")]
static mut PIDS: HashMap<u32, u32> = HashMap::<u32, u32>::with_ma

#[fentry(name = "kernel_clone")]
pub fn kernel_clone(ctx: FEntryContext) -> u32 {
    match try_kernel_clone(ctx) {
        Ok(ret) => ret,
        Err(_) => 1,
    }
}

fn try_kernel_clone(ctx: FEntryContext) -> Result<u32, c_long> {
    // Get the pid
    let pid = ctx.pid();
    // Save the pid in map.
    unsafe { PIDS.insert(&pid, &0, 0)? };
    Ok(0)
}
```

RUST TOOLCHAIN IS ALL YOU NEED

- No need for installing libbpf, make, clang or any C libraries in your system.

AYA-LOG

Logging library, based on perf buffers, production-ready.

```
#[fentry(name = "kernel_clone")]
pub fn kernel_clone(ctx: FEntryContext) -> u32 {
    let pid = ctx.pid();
    info!(&ctx, "new process: pid: {}", pid);
    0
}
```


ASYNC SUPPORT

- Userspace part of projects can be asynchronous.
- Support for Tokio and async-std.
- aya-template uses Tokio by default.

HOW TO START WITH AYA?

```
rustup toolchain install nightly --component rust-src  
cargo install bpf-linker  
cargo install cargo-generate
```

AYA-TEMPLATE

Creating a new project:

```
cargo install cargo-generate  
cargo generate https://github.com/aya-rs/aya-template
```



```
vadorovsky ~/playground > cargo generate https://github.com/aya-rs/aya-template
△ Unable to load config file: /home/vadorovsky/.cargo/cargo-generate.toml
△ Favorite https://github.com/aya-rs/aya-template not found in config, using it as a git repo url
🔑 Project Name : firewall
🔑 Generating template ...
? 🤖 Which type of eBPF program?
kprobe
kretprobe
fentry
fexit
uprobe
uretprobe
sock_ops
socket_filter
sk_msg
xdp
> classifier
cgroup_skb
cgroup_sysctl
cgroup_sockopt
tracepoint
lsm
tp_btf
```

PROJECT LAYOUT

```
— Cargo.toml
— firewall
  └─ Cargo.toml
    └─ src
      └─ main.rs
— firewall-common
  └─ Cargo.toml
    └─ src
      └─ lib.rs
— firewall-ebpf
  └─ Cargo.toml
  └─ rust-toolchain.toml
  └─ src
    └─ main.rs
— README.md
— xtask
  └─ Cargo.toml
  └─ src
    └─ build_ebpf.rs
    └─ main.rs
    └─ run.rs
```

RUNNING THE PROJECT

```
cargo xtask run
```

Only building:

```
cargo xtask build-ebpf  
cargo build
```

WHAT'S MISSING?

- Support for CO-RE in eBPF part of Aya
- That will require more work in Rust compiler

LET'S SEE A SIMPLE FIREWALL IN ACTION!

THANK YOU

- aya-rs.dev
- github.com/aya-rs/aya
- [Discord](#)